

Notes on R

Created by Brent Wagenseller, a GT Student

[Comments](#)

[Running R](#)

[Installing Packages](#)

[Saving screen output to a file](#)

[Vectors](#)

[Lists](#)

[Arrays](#)

[Saving / loading your current session](#)

[Importing data into data frames from CSVs](#)

[Importing data from SPSS files](#)

[Importing data from MySQL](#)

[Importing data from generic ODBC SQL](#)

[Exploring and understanding data](#)

[Boxplots / box-and-whisker plots](#)

[Histograms](#)

[Variance / Standard Deviation](#)

[Examining categorical variables](#)

Comments

"

multiline comments aren't natively supported in R - instead, wrap the entire line in double quotes like this. R will actually read the comment BUT it won't do anything with it

"

#A single-line comment is written like this - with a # at the beginning of the comment

#####Running R

"

Running R

Ubuntu

To run R in Ubuntu (after installation), simply type this at a prompt (capitalization is important): R

To quit, type: q()

Windows

For Windows, when R is installed a program 'RGui' is also installed - simply open that program. I assume you can type 'R' in at a command line on windows as well but I am not sure.

To quit R, type this at the R command prompt: q()

"

```
#####Running R
```

Scripts

#to run a script in R, note its path and on a command line run
source("my/pat/to/the/R_file.R")

Installing Packages

```
#####Installing packages
```

#Sometimes you need to load functions from a remote mirror / repository (R will list mirrors for you if need be)

#Note that you either have to be root for this (in UNIX) OR have to launch the current window as 'administrator' (in Windows)

#OR be OK with R storing the package locally (so all users will have to download the package if they wish to use it); R give you the choice when you load a package / library

#this installs the "foreign" package, which allows for the importation of SPSS data in its native format

#only has to be done once (as root; if installed as a user it must be installed once per user)

#Note the double quotes may not be necessary (sometimes they are, sometimes they arent) but Google recommends them in their R standards

```
install.packages("foreign")
```

#this is the way to load a package / library from your local machine (once installed)

#this loads the example "foreign" library

```
library("foreign")
```

```
#This is used for when you actually create a function or package – this is used instead of
library("foreign")
#This can be used in place of library("foreign") for the command line too, but most people just
use library("foreign")
require("foreign")
```

Unload Package

```
#This unloads – but does not delete – the package 'foreign'
#By default, all installed packages are removed when R quits
detach("package:foreign", unload=TRUE)
```

```
#Checks for updates; this is unfortunately a manual thing. Do this periodically!
update.packages()
#####Timing Code
#create a start time variable
startTime <- proc.time()
proc.time() - startTime#displays to screen; elapsed time is the time it took to run

#conversely, you can run code inside system.time() and the output is the time
system.time(a <- c(1,2,3))
```

Saving screen output to a file

```
#get working directory
getwd()

#set working directory
setwd('/home/myHomeDirectory')

#open a file for writing
sink("sink.txt")
```

```
#modify a variable
```

```
c <- c(1,2,3,4,5)
```

```
#print to file (remember, output goes to file and not the screen)
```

```
c
```

```
#return the output to the screen
```

```
sink()
```

```
#delete the file
```

```
unlink('sink.txt')
```

Vectors

"

The fundamental R data structure is the vector, which stores an ordered set of values called elements.

Vectors can only contain the same type - numbers OR character strings but not both.

Vector types:

integer

numeric(decimal)

character

logical (true or false)

Additionally any vector can have:

NULL (absence of any value)

NA (missing value)

"

```
#Vector examples:
```

```
#note: the c() function is read 'combine' and combines elements to make a vector
```

```
#here, the vector 'subject_name' is created and three values are stored to it
```

```
#note the use of '<-' - this is used in place of '=' for assignment in R for the most part
```

```
#also note that simply writing a variable name on a line in R will print it to the screen
```

```
subjectName <- c("John", "Jane", "Steve")
```

```
#to reference specific values in the vector you can reference them by a range of numbers  
corresponding to their position like so
```

```
subjectName[2:3]
```

```
#you can exclude certain positions with a minus sign:
```

```
subjectName[-2]
```

```
#you can also construct a logical vector on the fly to tell R which elements to include
subjectName[c(FALSE, TRUE, TRUE)]
```

```
#'nominal' data is features that represent a characteristic with categories of values (for example,
cars:ford/honda, meats:chicken/veal/fish/beef, etc - basically anything not numbers)
```

```
#R provides a data structure known as a 'factor' to store nominal data
```

```
#its important to use factors for string-based categories as R can efficiently store them with an
internal ID; if a factor isnt used it has to save the actual string for each row, also some
algorithms REQUIRE nominals to be used in factors
```

```
#Example
```

```
gender <- factor(c("MALE", "FEMALE", "MALE"))
```

```
#factors are also useful if there is another data type that is not represented in our data - yet - but
we want to acknowledge exists. For example:
```

```
bloodType <- factor(c("O", "AB", "A"), levels = c("A", "B", "AB", "O"))
```

```
#note that 'B' is not used but still noted as a possible type
```

Lists

```
#Lists are similar to vectors but do NOT require each element to be of the same type (so
character strings and numbers can coexist in a list)
```

```
#Lists are used often because of this
```

```
#as an example, lets create a list - subject1 - and populate it with some of the variables we
already created
```

```
subject1 <- list(fullname = subjectName[1], gender = gender[1], bloodType = bloodType[1])
```

```
#subject1 is now a list! be warned though: this is an actual list; in this instance we just use one
record but a list can hold several records
```

```
#note that if we wanted to get a specific column in the list we just created, its the variable and
column separated by a $ like so
```

```
subject1$fullname
```

```
#####Data Frames
```

```
#this is by far the most important data structure in R
```

```
#similar to a spreadsheet or database since it has rows and columns
```

```
#a data frame is like a list of vectors or factors, each having exactly the same number of
values
```

```
#lets create a data frame
```

```
#note: patient.data is a variable name below - its NOT a function or an object. the '.' is not
treated as an object in R variables but simply another character
```

#note the use of 'stringAsFactors = FALSE' - this tells R to NOT convert strings to factors; mainly because the name is definitely not a 'category' and should not be treated as such, you CAN go back later and specify c
patient.data <- data.frame(subjectName, gender, bloodType, stringAsFactors = FALSE)

#you can eliminate a specific row by using a construct and listing only the columns/vectors/frames you want to list
#whats cool about this is information is extracted by column name
patient.data[c("subjectName", "bloodType")]

#you can treat the data frame much like a multi-dimensional array in other languages; use the row number then the column number
#here is how you can extract the gender of the third person
patient.data[3,2]

#####Matrixes

#A matrix is a data structure that represents a two dimensional table
#most often used for mathematical operations so they usually only store numbers, but they can store characters as well

#creates a matrix of a/b/c/d with two rows; how this will work is the first row will be a and b and the second row will be c and d
#note that R loads the data into a COLUMN first, NOT across rows
myMatrix <- matrix(c('a', 'b','c','d'),nrow=2)

Arrays

#apparently very similar to a matrix yet while a matrix just has rows and columns, an array has rows, columns, and several other layers
#arrays are not used often and are apparently somewhat advanced - so they are mentioned but not used in my training documents

Saving / loading your current session

#since R usually isn't run by scripts sequentially - and is really command-line driven - there is a method for saving variables for a later time; to save specific data frame / matrices / arrays / lists / frames / etc:
save(patient.data, myMatrix, file = 'learnR.RData')
#note the variables were listed first, then a file was assigned

#to load this data later back from that file:
load("learnR.RData")

#also note the save.image() command saves the ENTIRE session to a file called .RData
#By default, R looks for this file when loading initially to restore the previous session

Importing data into data frames from CSVs

#While possible to import excel files, they are special packages and they do not always work;
instead, use CSV

#R converts missing data to "NA"

#For windows, the normal backslashes for directories will not work; instead you either have to
use double backslashes OR use front slashes a la UNIX

#read in a csv file into the dataframe 'sn.csv' and signify that the first row contains column
headers (otherwise R names them V1 V2 V3 etc etc)

#note the 'sn.csv' is NOT a function call to csv; instead, its a fancy way to name a variable; the
entire variable name IS 'sn.csv' – dots can be used in variable names and Google suggests it to
separate words in variable names

#On the other hand, 'read.csv' is the function that reads the CSV file into the variable
sn.csv <- read.csv("/home/me/social_network.csv", header = T)

#this is the 'structure' function, and it transposes the columns and rows and then prints them to
the screen

str(sn.csv)

#write a data frame to a csv file

write.csv(patient.data, file = 'patientData.csv')

Importing data from SPSS files

#Very common statistical package used by researchers

#SPSS is very much like an Excel spreadsheet
#In SPSS, its common to see a string column immediately followed by a number column that represents the string column (redundant)
#for example, there may be a column that is strictly limited to male/female and then the next column would be a 1(for male) or 0 (for female).
#For the column that is a number, you can additionally place a 'label' for each numeric value (separate from the associated string column)
#A program that can edit SPSS files is 'IBM SPSS Statistics Data Editor'
#Lynda suggests saving it as a CSV file and then importing it using the command 'read.csv'
sn.spss.f <- read.spss("/home/me/social_network.sav", to.data.frame=T, use.value.labels=T)
#to.data.frame=T means we want to convert this to a data.frame (not sure why you wouldnt do this)
#use.value.labels=T means use the associated labels with all numeric columns that have labels

Importing data from MySQL

#Its possible to interact with a MySQL Database in R
#more info at: www.r-bloggers.com/mysql-and-r/
#Installation in Ubuntu - Before you run the install command in R, you need to install a package in Ubuntu itself (as root) via: apt-get install libmysqlclient-dev
#Installation in Windows - when I did this in windows I simply loaded the package (no extra steps required)

#In R itself: After following the steps above for your platform, run the following in R:
install.packages("RMySQL")

#load the MySQL library
library("RMySQL")

#this is the connection string
#user/password – self-explanatory
#host = the IP (or hostname) of the MySQL database
#dbname – the database name / schema name
con <- dbConnect(MySQL(), user = 'userName', password = 'userPassword', host = '127.0.0.1', dbname='mySchemaName')

#This writes a data frame to a table
#conn – the connection string
#name – the name of the table. Note if the schema is different than the one supplied on the connection string you can reference it here as you would in MySQL a la Schema.TableName
#value – the data frame variable


```
dbWriteTable(conn=con, name = 'Test2', value = patient.data)
```

```
#This loads a tables contents into a data frame ('fuzzy' is the data frame that will hold the new data)
```

```
#conn – the connection string
```

```
#name – the name of the table. Note if the schema is different than the one supplied on the connection string you can reference it here as you would in MySQL a la Schema.TableName.
```

```
#Note: Views can be used – so if you want to massage data in a table, simply create a view and use that in place of the table in R!
```

```
fuzzy = dbReadTable(conn=con,name='Test2')
```

```
#if you wish to write a query, this is how it's done
```

```
#preps the 'sql' variable for a SQL query; note varId and varName are not set in this example but they should be set prior to this point
```

```
sql <- sprintf("SELECT * FROM MySchema.MyTable WHERE ID = %d AND myName = '%s'",  
varId, varName)
```

```
#sets the data frame 'rs' to the query results
```

```
rs <- dbGetQuery(con, sql)
```

```
#disconnects the database
```

```
dbDisconnect(con)
```

Importing data from generic ODBC SQL

```
#This will work for Oracle, MySQL, Microsfot SQL, SQLite, or PostgreSQL (possibly HP Vertica)
```

```
#The instructional book didnt get into connection strings - which are needed. however the connection string may be
```

```
#able to be pulled from a perl/python/php script (I havent tried it).
```

```
#the package 'RODBC' is needed
```

```
install.packages("RODBC")
```

```
library("RODBC")
```

```
#save your connection string here
```

```
my_dsn <- "my awesome connection string"
```

```
#create a connection
```

```
myGenericConnection <- odbcConnect("my_dsn")
```

```
#IF the connection requires a separate username and password, use this method instead
myGenericConnection <- odbcConnect("my_dsn", uid = "myUserName", pwd = "myPAssword")

#create the query and pull into data frame 'newPatientDataFrame'
patientQuery <- "SELECT * FROM MyTable WHERE alive = 1"
newPatientDataFrame <- sqlQuery(channel = myGenericConnection, query = patientQuery,
stringAsFactors = FALSE)

#close connection
odbcClose(myGenericConnection)
```

Exploring and understanding data

```
#This section coincides with the lesson in 'Machine Learning in R' by Brett Lantz, section
'Exploring the structure of data'
#on p43. The whole point of R is data analysis; the code is important, but knowing how to USE
R is also important
```

```
#this section explores data surrounding used cars
#pull in the data to a data frame
usedcars <- read.csv("usedcars.csv", stringsAsFactors = FALSE)
```

```
#A good first start is looking at the summary statistics of the data - the str() (structure) function is
good at getting a first look
str(usedcars)
```

```
#We note that this seems to be data on used cars - but we need to know more
#We pick a variable - year - and explore it using the 'summary;' command
summary(usedcars$year)
#Year could be deceiving - is it the year the car was made or the year the car was sold to the
dealership?
#the summary data shows the minimum is 2000, so its a safe bet its the year the car was made
```

```
#now take a look at the summary of two more columns - price and mileage
summary(usedcars[c("price", "mileage")])
```

```
#now lets look at some other numbers. Again, read chapter 2 of 'Machine Learning with R' by
Brett Lantz to get a
#better understanding of how to use this data
```

```
# the min/max of used car prices
range(usedcars$price)
```

```
# the difference of the range
diff(range(usedcars$price))
```

```
# IQR for used car prices; IQR is the 'interquartile range', which is Q3 - Q1. This is helpful to
determine the spread of data
IQR(usedcars$price)
```

```
# use quantile to calculate five-number summary; without parameters set, this returns the 5
basic quantiles (that is, 0/25/50/75/100)
quantile(usedcars$price)
```

```
# get the 1st and the 99th percentile
quantile(usedcars$price, probs = c(0.01, 0.99))
```

```
# If we wanted to find the quintiles of the data (cutting it into 5 groups with an equal number of
data points per group)
#we can use the seq() (sequence function) which generates vectors of evenly-spaced values
quantile(usedcars$price, seq(from = 0, to = 1, by = 0.20))
```

Boxplots / box-and-whisker plots

```
#This is how a boxplot is created in R
#Boxplots are useful in determining how clustered the data is
#Boxplots can be highly customized - type ?boxplot for more info
# boxplot of used car prices and mileage
boxplot(usedcars$price, main="Boxplot of Used Car Prices", ylab="Price ($)")
boxplot(usedcars$mileage, main="Boxplot of Used Car Mileage", ylab="Odometer (mi.)")
```

Histograms

```
#This is how a boxplot is created in R
#Histograms are useful in determining how clustered the data is
#unlike boxplots, histograms do not divide the data into 4 (for quartile) equal bins; instead, they
use a constant bin size and report how many variables / data points fall in that bin
#Histograms can be highly customized - type ?histograms for more info
# histograms of used car prices and mileage
```

```
hist(usedcars$price, main = "Histogram of Used Car Prices", xlab = "Price ($)")  
hist(usedcars$mileage, main = "Histogram of Used Car Mileage", xlab = "Odometer (mi.)")
```

Variance / Standard Deviation

```
# variance and standard deviation of the used car data  
#see the notes on variance and standard deviation in chapter 2 of 'Machine Learning with R'  
  
#note that R uses the sample variance (n-1) and NOT the typical population variance (n) in the  
variance and standard deviation equations.  
#This shouldnt effect the outcome too much provided the dataset isnt too small (5-1 has much  
more of an impact than 500-1)
```

```
var(usedcars$price)  
sd(usedcars$price)  
var(usedcars$mileage)  
sd(usedcars$mileage)
```

Examining categorical variables

```
# one-way tables for the used car data - tells us the count of each category  
#a 'table' is not like a table in a database - rather, its a way to count the different number of  
occurrences in a categorical variable data set  
table(usedcars$year)  
table(usedcars$model)  
table(usedcars$color)  
  
# compute table proportions - what % of the table represents Toyota?  
model_table <- table(usedcars$model)  
prop.table(model_table)  
  
# make the above data more recognizeable as a percentage by multiplying it by 100, then round  
the data to 1 decimal
```

```
color_table <- table(usedcars$color)
color_pct <- prop.table(color_table) * 100
round(color_pct, digits = 1)
```

```
# scatterplot of price vs. mileage
plot(x = usedcars$mileage, y = usedcars$price,
     main = "Scatterplot of Price vs. Mileage",
     xlab = "Used Car Odometer (mi.)",
     ylab = "Used Car Price ($)")
```

```
# new variable indicating conservative colors
usedcars$conservative <-
usedcars$color %in% c("Black", "Gray", "Silver", "White")
```

```
# checking our variable
table(usedcars$conservative)
```

```
library("gmodels")
# Crosstab of conservative by model
#chisq = TRUE means the chi-squared test for independence is included; its in the last line of
output, p = number.
#in this case p = .92591; since its high it means there probably isnt a correlation between model
and conservative colors
CrossTable(x = usedcars$model, y = usedcars$conservative, chisq = TRUE)
```